

Advanced OpenMP

OpenMP Tips, tricks and gotchas



Directives

- Mistyping the sentinel (e.g. `!OMP` or `#pragma opm`) typically raises no error message.
 - Be careful!
 - Extra nasty if it is e.g. `#pragma opm atomic` – race condition!
 - Write a script to search your code for your common typos

Writing code that works without OpenMP too

- The macro **`_OPENMP`** is defined if code is compiled with the OpenMP switch.
 - You can use this to conditionally compile code so that it works with and without OpenMP enabled.
- If you want to link dummy OpenMP library routines into sequential code, there is code in the standard you can copy (Appendix A in 4.0)

Parallel regions

- The overhead of executing a parallel region is typically in the tens of microseconds range
 - depends on compiler, hardware, no. of threads
- The sequential execution time of a section of code has to be several times this to make it worthwhile parallelising.
- If a code section is only sometimes long enough, use the **if** clause to decide at runtime whether to go parallel or not.
 - Overhead on one thread is typically much smaller ($<1\mu\text{s}$).
- You can use the EPCC OpenMP microbenchmarks to do detailed measurements of overheads on your system.
- Download from
www.epcc.ed.ac.uk/research/computing/performance-characterisation-and-benchmarking

Is my loop parallelisable?

- Quick and dirty test for whether the iterations of a loop are independent.
- Run the loop in reverse order!!
- Not infallible, but counterexamples are quite hard to construct.

Loops and nowait

```
#pragma omp parallel
{
#pragma omp for schedule(static) nowait
    for(i=0;i<N;i++){
        a[i] = ....
    }
#pragma omp for schedule(static)
    for(i=0;i<N;i++){
        ... = a[i]
    }
}
```

- This is safe so long as the number of iterations in the two loops and the schedules are the same (must be static, but you can specify a chunksize)
- Guaranteed to get same mapping of iterations to threads.

Default schedule

- Note that the default schedule for loops with no schedule clause is implementation defined.
- Doesn't have to be STATIC.
- In practice, in all implementations I know of, it is.
- Nevertheless you should not rely on this!
- Also note that SCHEDULE(STATIC) does not completely specify the distribution of loop iterations.
 - don't write code that relies on a particular mapping of iterations to threads

Tuning the chunksize

- Tuning the chunksize for static or dynamic schedules can be tricky because the optimal chunksize can depend quite strongly on the number of threads.
- It's often more robust to tune the *number of chunks per thread* and derive the chunksize from that.
 - chunksize expression does not have to be a compile-time constant

SINGLE or MASTER?

- Both constructs cause a code block to be executed by one thread only, while the others skip it: which should you use?
- MASTER has lower overhead (it's just a test, whereas SINGLE requires some synchronisation).
- But beware that MASTER has no implied barrier!
- If you expect some threads to arrive before others, use SINGLE, otherwise use MASTER

Data sharing attributes

- Don't forget that private variables are uninitialised on entry to parallel regions!
- Can use **firstprivate**, but it's more likely to be an error.
 - use cases for firstprivate are surprisingly rare.

Default(none)

- The default behaviour for parallel regions and worksharing construct is **default(shared)**
- This is extremely dangerous - makes it far too easily to accidentally share variables.
- Possibly the worst design decision in the history of OpenMP!
- Always, always use **default(none)**
 - I mean always. No exceptions!
 - Everybody suffers from “variable blindness”.

Spot the bug!

```
#pragma omp parallel for private(temp)
  for(i=0;i<N;i++){
    for (j=0;j<M;j++){
      temp = b[i]*c[j];
      a[i][j] = temp * temp + d[i];
    }
  }
```

- May always get the right result with sufficient compiler optimisation!

Private global variables

```
double foo;
```

```
#pragma omp parallel \  
private(foo)  
{  
    foo = ....  
    a = somefunc();  
}
```

```
extern double foo;  
  
double sumfunc(void) {  
    ... = foo;  
}
```



- Unspecified whether the reference to **foo** in **somefunc** is to the original storage or the private copy.
- Unportable and therefore unusable!
- If you want access to the private copy, pass it through the argument list (or use **threadprivate**).

Huge long loops

- What should I do in this situation? (typical old-fashioned Fortran style)

```
do i=1,n  
..... several pages of code referencing 100+  
          variables  
end do
```


- Determining the correct scope (private/shared/reduction) for all those variables is tedious, error prone and difficult to test adequately.

- Refactor sequential code to

```
do i=1,n  
    call loopbody(.....)  
end do
```

- Make all loop temporary variables local to loopbody
- Pass the rest through argument list
- Much easier to test for correctness!
- Then parallelise.....
- C/C++ programmers can declare temporaries in the scope of the loop body.

Reduction race trap

```
#pragma omp parallel shared(sum, b)
{
    sum = 0.0; 
    #pragma omp for reduction(+:sum)
    for(i=0;i<n;i++) {
        sum += b[i];
    }
    .... = sum;
}
```

- There is a race between the initialisation of **sum** and the updates to it at the end of the loop.

Missing SAVE or static

- Compiling my sequential code with the OpenMP flag caused it to break: what happened?
- You may have a bug in your code which is assuming that the contents of a local variable are preserved between function calls.
 - compiling with OpenMP flag forces all local variables to be **stack allocated** and not heap allocated
 - might also cause stack overflow
- Need to use SAVE or static correctly
 - but these variables are then shared by default
 - may need to make them threadprivate
 - “first time through” code may need refactoring (e.g. execute it before the parallel region)

Stack size

- If you have large private data structures, it is possible to run out of stack space.
- The size of thread stack *apart from the master thread* can be controlled by the **OMP_STACKSIZE** environment variable.
- The size of the master thread's stack is controlled in the same way as for sequential program (e.g. compiler switch or using **ulimit**).
 - OpenMP can't control this as by the time the runtime is called it's too late!

Critical and atomic

- You can't protect updates to shared variables in one place with atomic and another with critical, if they might contend.
- No mutual exclusion between these
 - critical protects code, atomic protects memory locations.

```
#pragma omp parallel
{
    #pragma omp critical
        a+=2;
    #pragma omp atomic
        a+=3;
}
```

Allocating storage based on number of threads

- Sometimes you want to allocate some storage whose size is determined by the number of threads.
 - but how do you know how many threads the next parallel region will use?
- Can call **omp_get_max_threads()** which returns the value of the *nthreads-var* ICV. The number of threads used for the next parallel region will not exceed this
 - except if a **num_threads** clause is used.
- Note that the implementation can always deliver fewer threads than this value
 - if your code depends on there actually being a certain number of threads, you should always call **omp_get_num_threads()** to check

Environment for performance

- There are some environment variables you should set to maximise performance.
 - don't rely on the defaults for these!

OMP_WAIT_POLICY=active

- Encourages idle threads to spin rather than sleep

OMP_DYNAMIC=false 

- Don't let the runtime deliver fewer threads than you asked for

OMP_PROC_BIND=true

- Prevents threads migrating between cores

Debugging tools

- Traditional debuggers such as DDT or Totalview have support for OpenMP
 - debugger changes the timing of event on different threads
- This is good, but they are not much help for tracking down race conditions
- Race detection tools work in a different way
 - capture all the memory accesses during a run, then analyse this data for races which *might have* occurred.
 - Intel Inspector XE

Timers

- Make sure your timer actually does measure wall clock time!
- Do use `omp_get_wtime()` !
- Don't use `clock()` for example
 - measures CPU time accumulated across all threads
 - no wonder you don't see any speedup.....

Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Acknowledge EPCC as follows: “© EPCC, The University of Edinburgh, www.epcc.ed.ac.uk”

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.