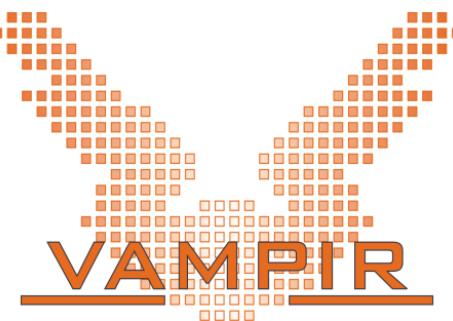


# Performance Analysis Exercises with Vampir

---

Matthias Weber, Holger Brunst, Hartmut Mix  
Technische Universität Dresden

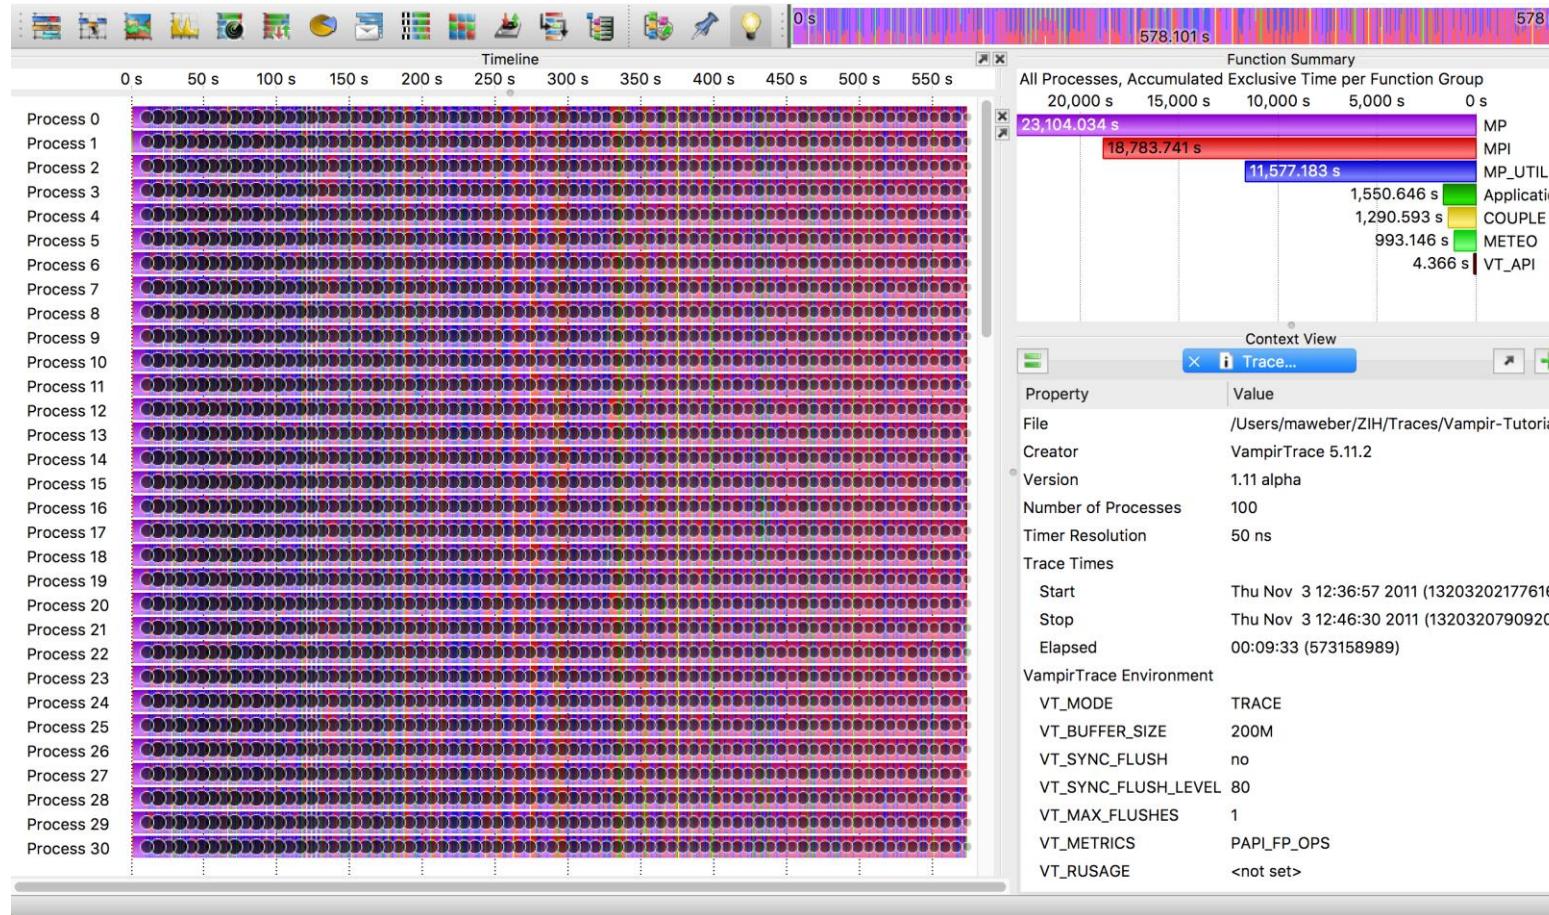


# Exercise Trace Files

```
% ls $TW35/trace-examples/\  
Vampir-Tutorial-Analysis-Examples/  
  
01-p100-cosmo-specs-orig  
02-p100-cosmo-specs-fd4  
03_wrf_deimos  
04_sbmfd4_jugene
```

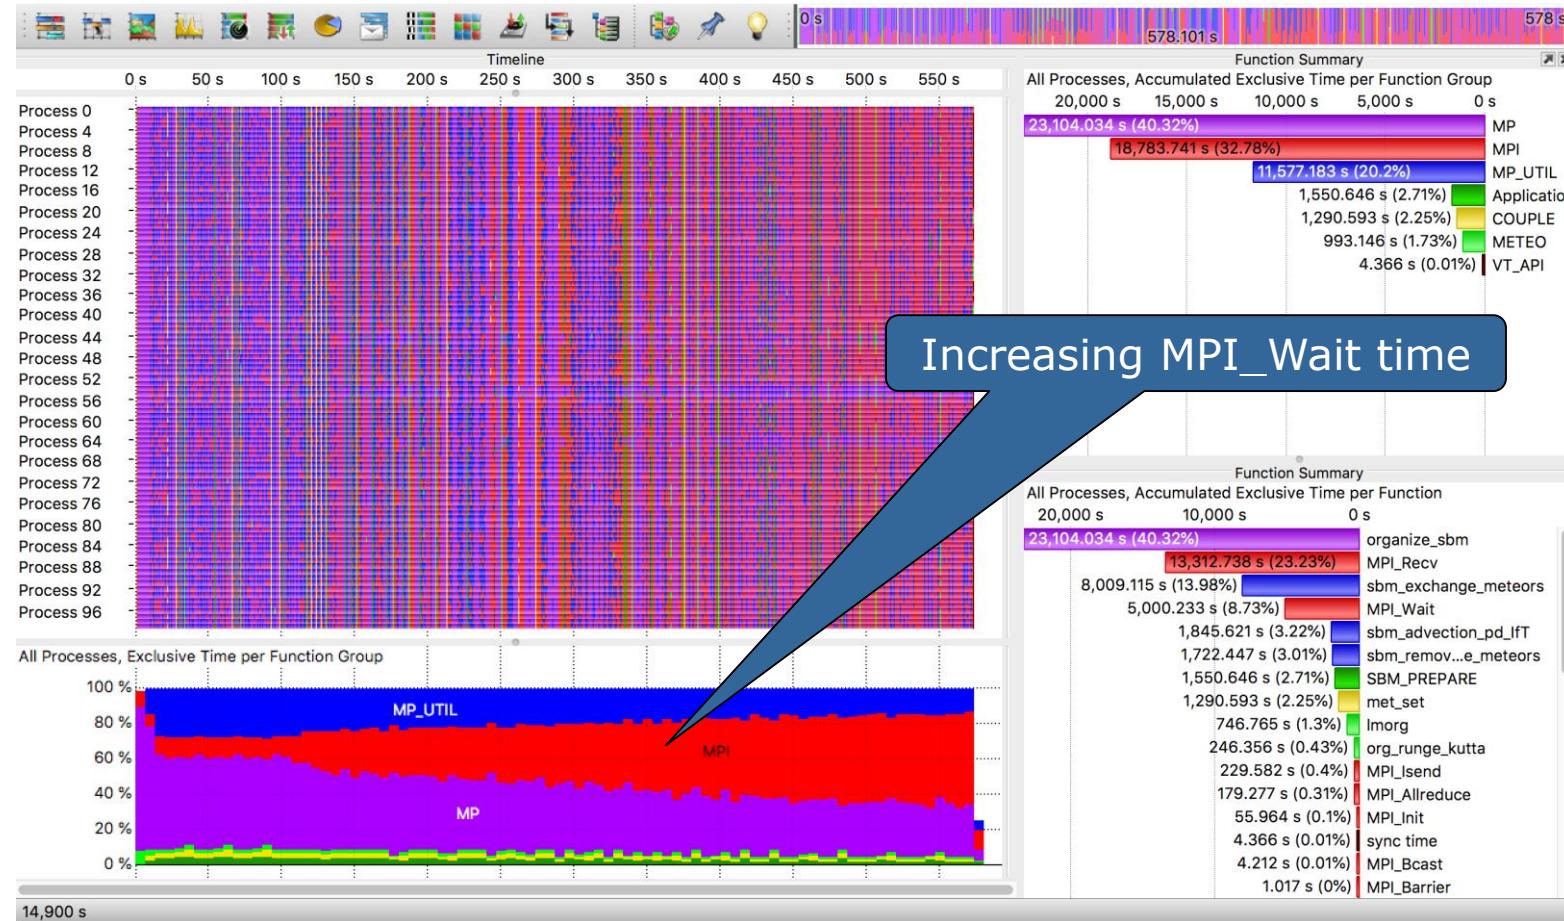
- Four trace files for exercising performance analysis with Vampir
- Traces show real application runs
- Do the traces contain performance problems?
- If yes, try to find their causes

# 01-p100-cosmo-specs-orig



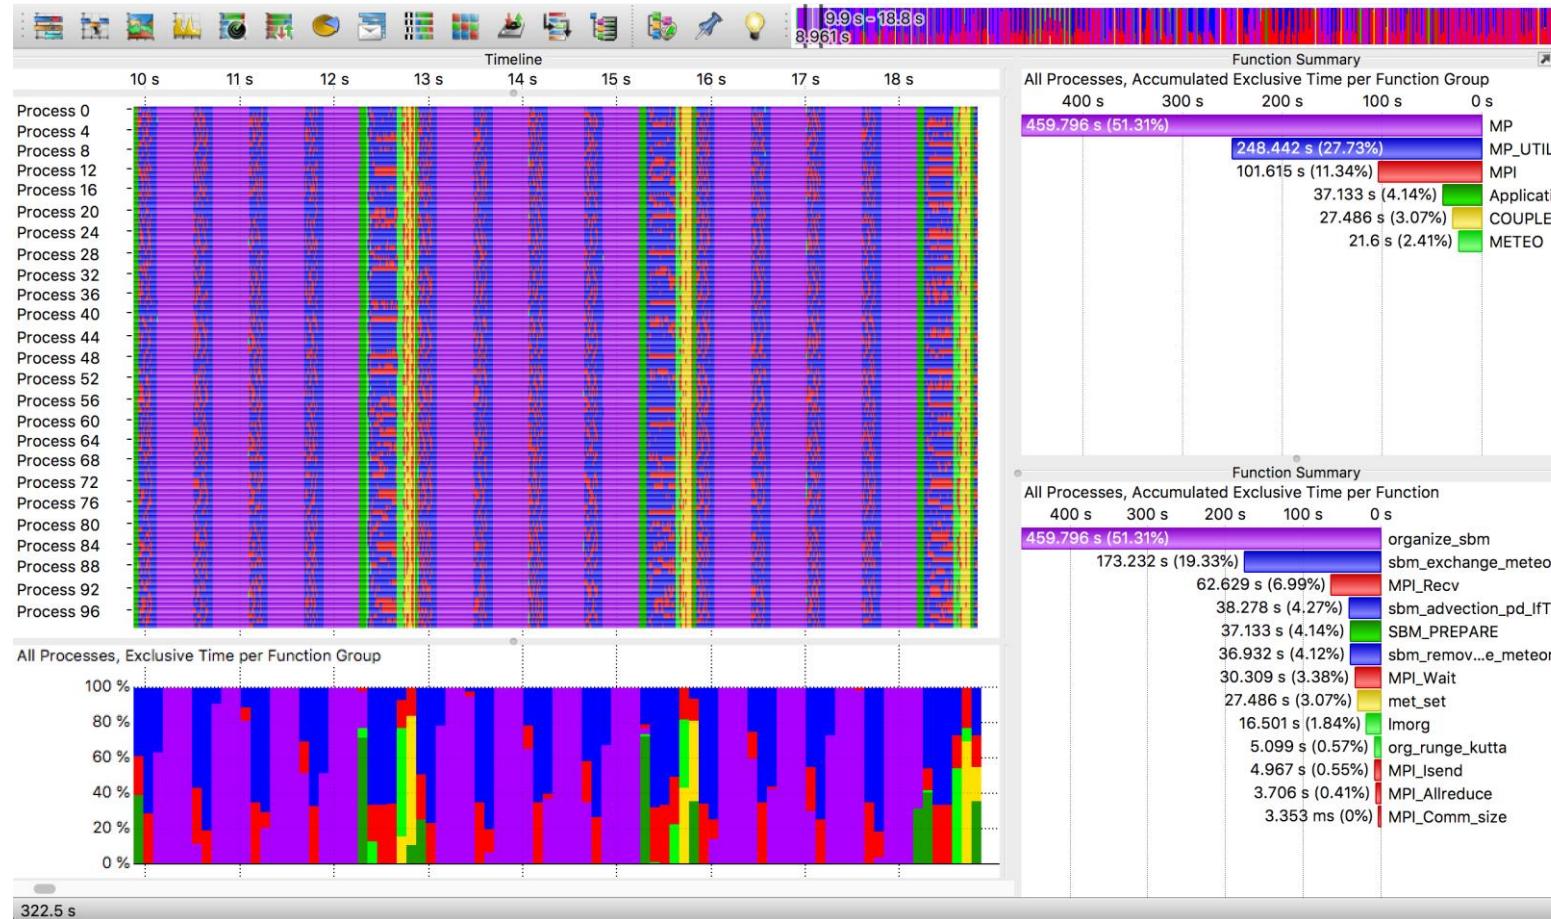
- Weather forecast code COSMO-SPECS
- Run with 100 processes
- COSMO: weather model (METEO group)
- SPECS: microphysics for accurate cloud calculation (MP and MP\_UTIL group)
- Coupling of both models done in COUPLE group

# 01-p100-cosmo-specs-orig



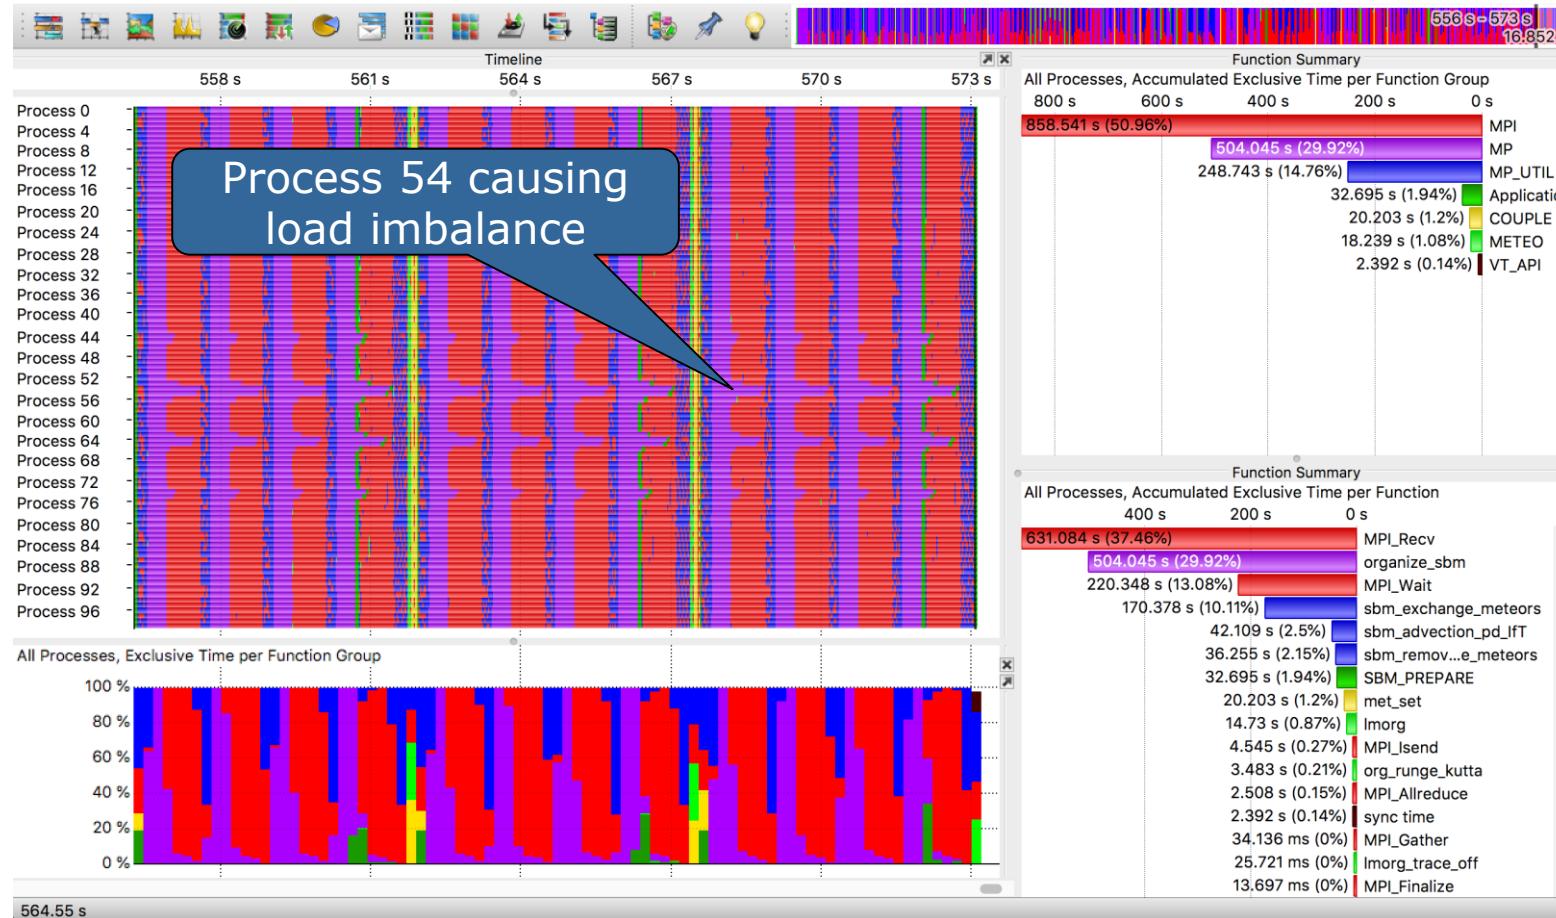
- Compared to METEO, MP and MP\_UTIL are very compute intensive, however this is due to more complex calculations and no performance issue
- Problem: >32% of time spent in MPI
- MPI runtime share increases throughout the application run

# 01-p100-cosmo-specs-orig



- Zoom into the first three iterations
- MP/MP\_UTIL perform four sub-steps in one iteration
- Low MPI time share
- Everything is balanced and looks okay

# 01-p100-cosmo-specs-orig



- Zoom into the last three iterations
- Very high MPI time share (>50%)
- Large load imbalance caused by MP functions around **Process 54** and **Process 64**

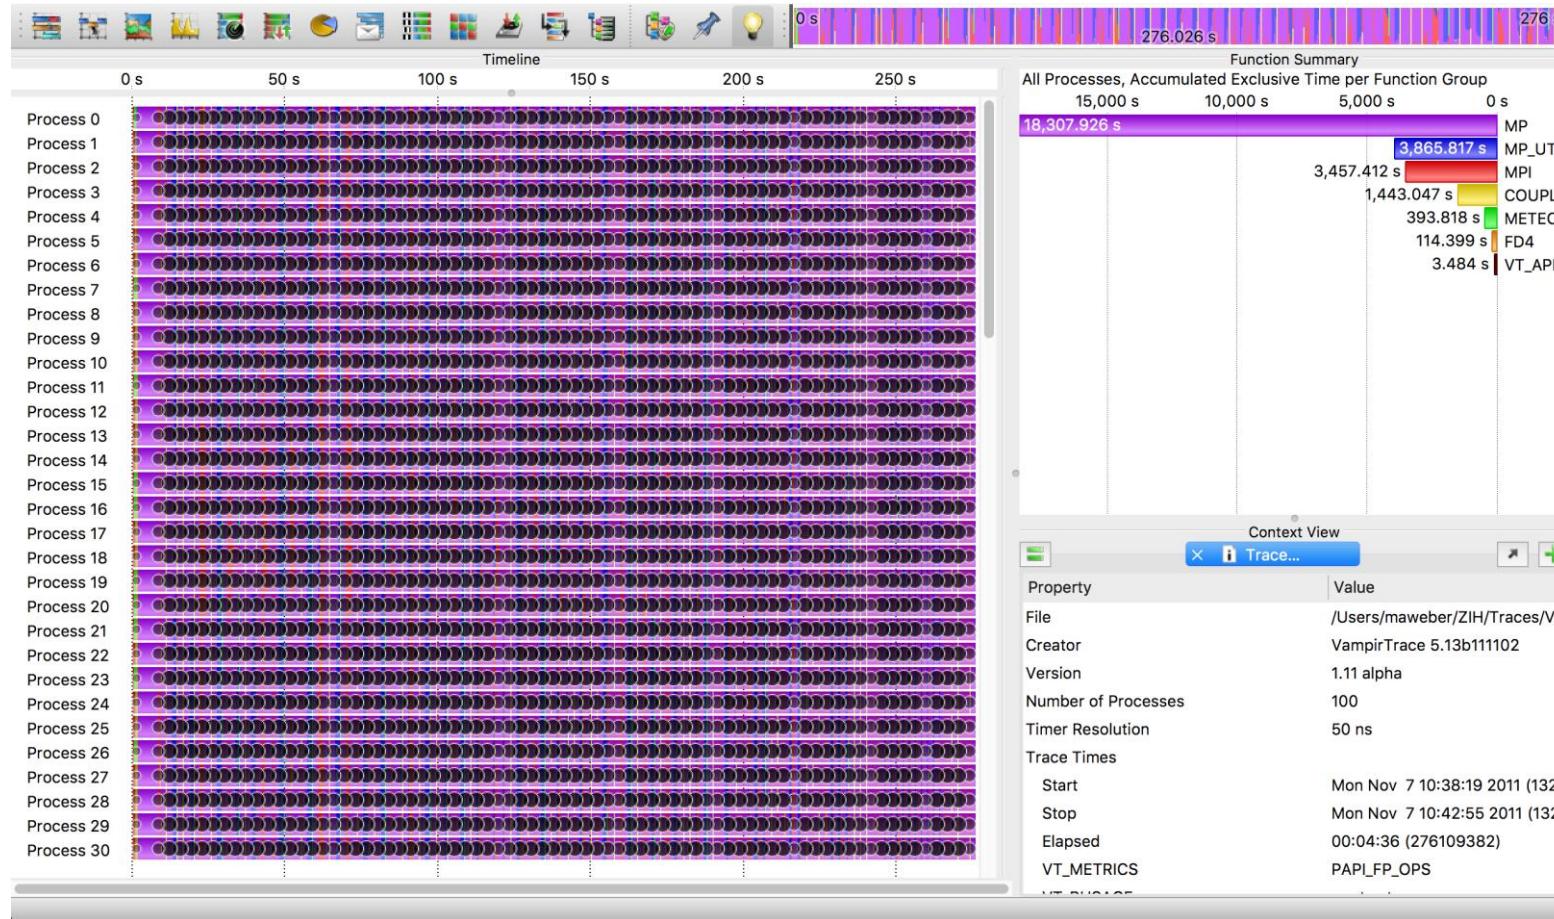
## 01-p100-cosmo-specs-orig

High FLOPs rates due to computation of clouds in this area



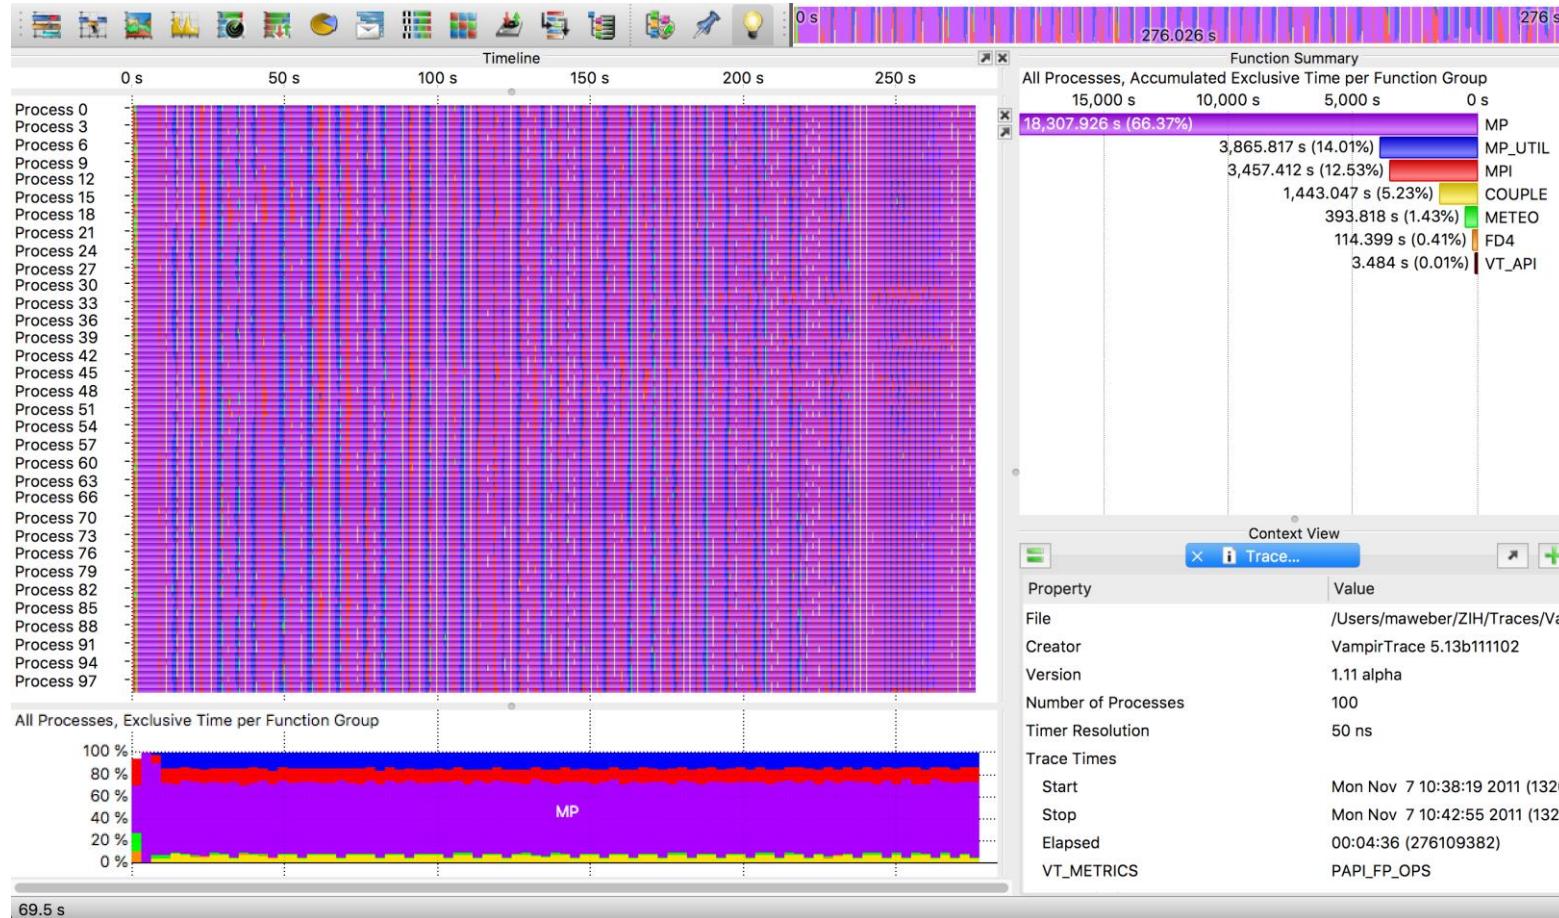
- **PAPI\_FP\_OPS** counter showing higher FLOPs rates on processes causing the imbalance
- Reason for imbalance:  
Static grid used for distribution of processes.  
Depending on the weather, expensive cloud computations (MP group) may be only necessary on some processes

## 02-p100-cosmo-specs-fd4



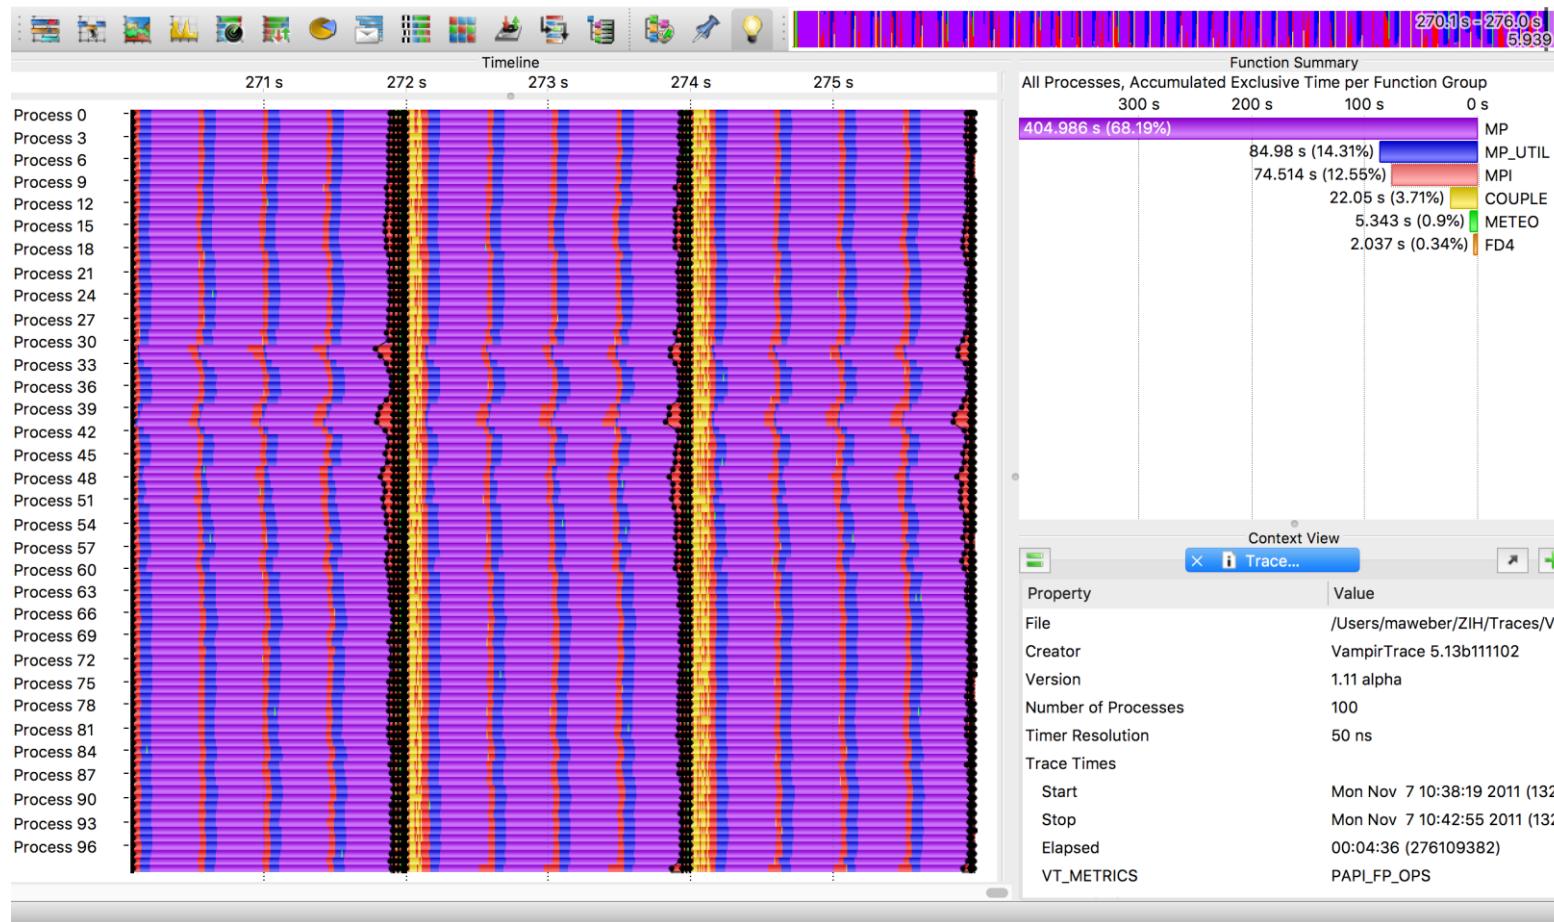
- Weather forecast code COSMO-SPECS
- Run with 100 processes
- COSMO: weather model (METEO group)
- SPECS: microphysics for accurate cloud calculation (MP and MP\_UTIL group)
- Coupling of both models done in COUPLE group
- Dynamic load balancing (FD4 group)

## 02-p100-cosmo-specs-fd4



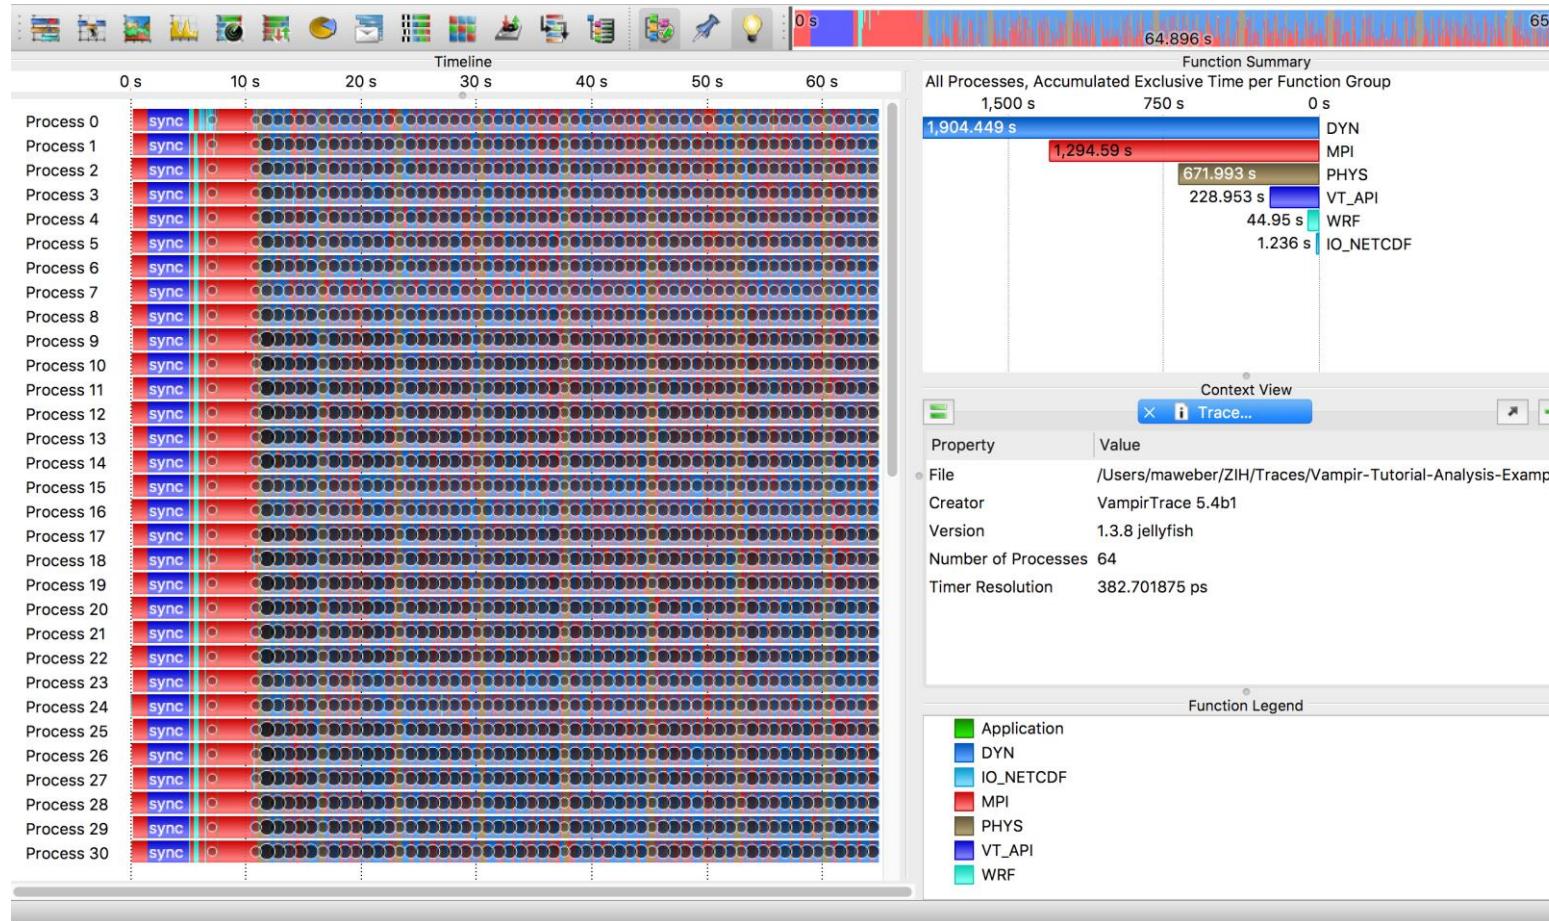
- Dynamic load balancing mitigates the balance problems of the original COSMO-SPECS version
- MPI time share is reduced to <13%
- MPI time share stays constant throughout the application runtime
- Runtime reduced by factor of 2.1, from initially 578s to 276s

## 02-p100-cosmo-specs-fd4



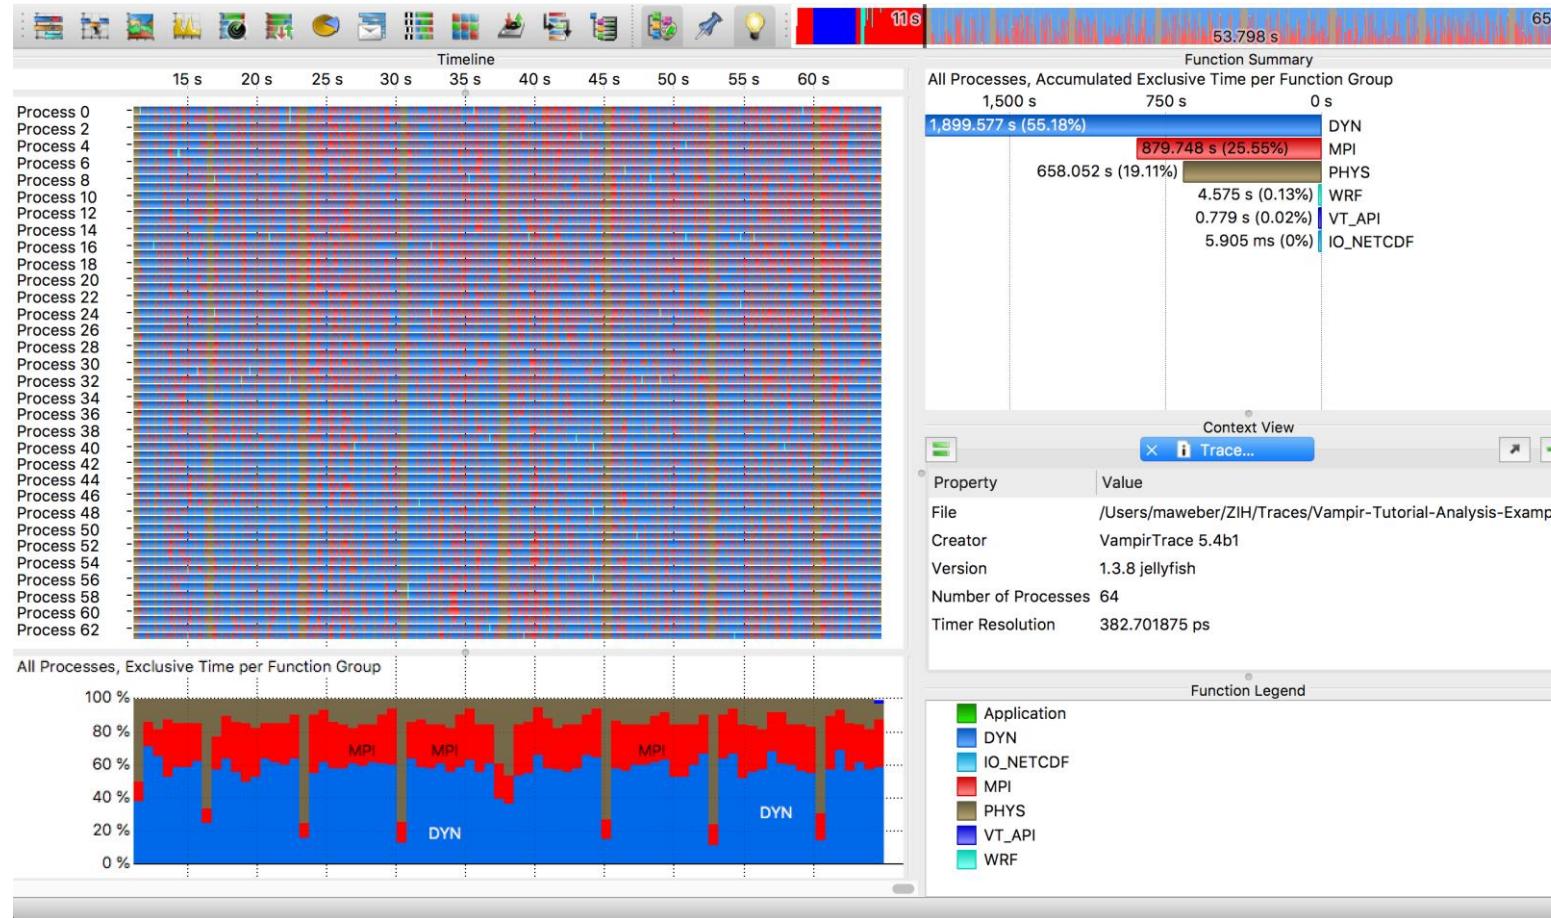
- Zoom into last three iterations
- FD4 balances MP load (precipitation processes in clouds) across all available processes

## 03\_wrf\_deimos



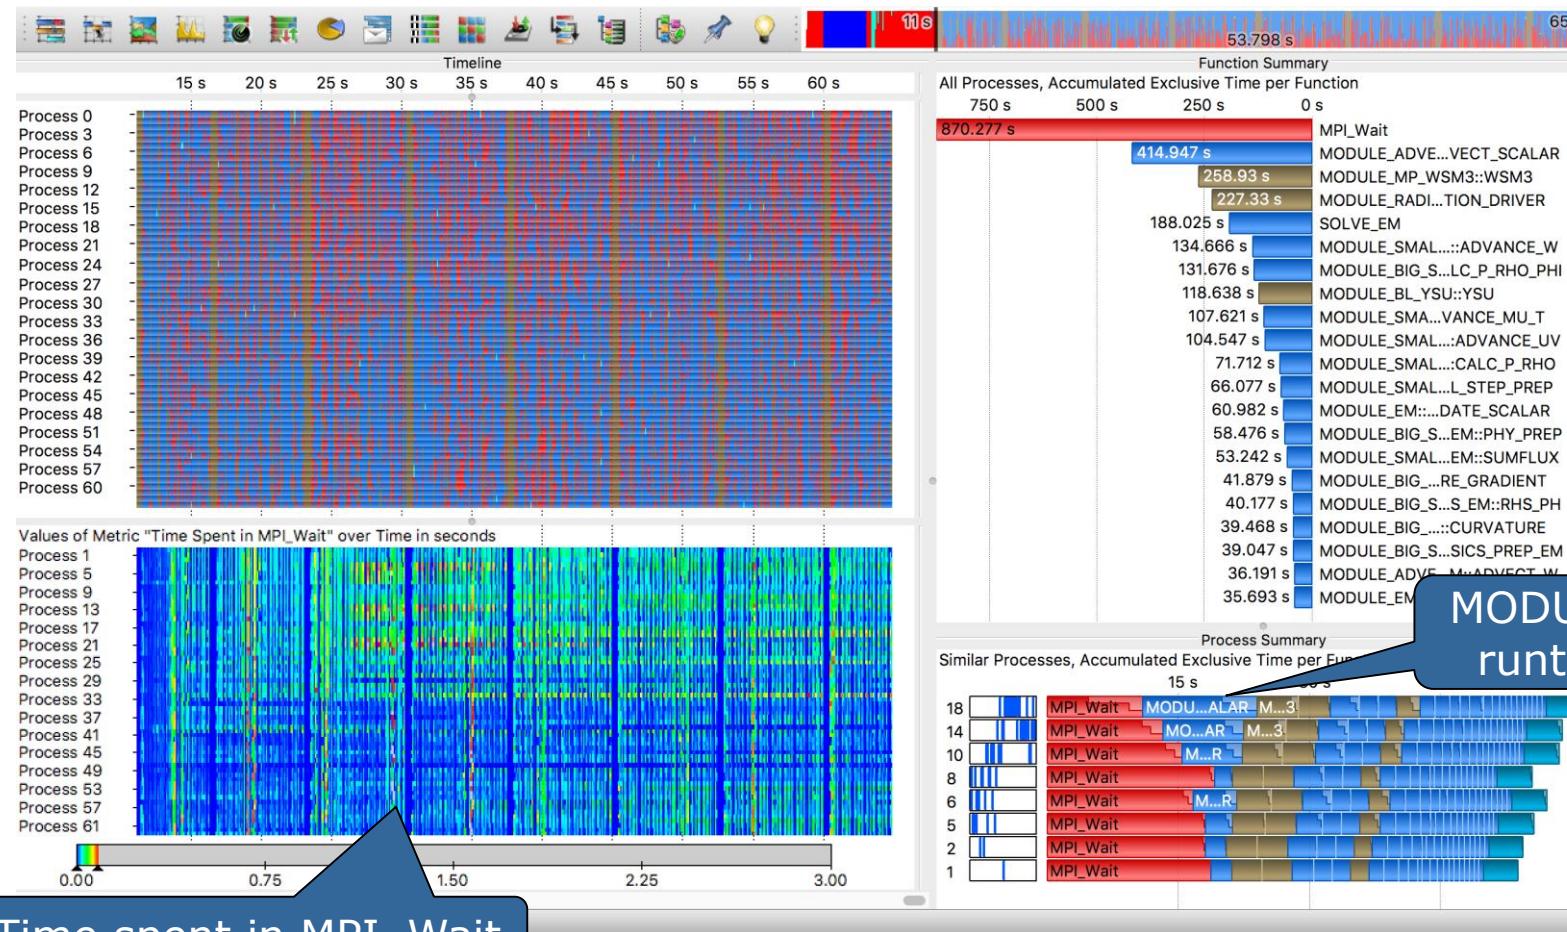
- Weather forecast code WRF
- Run with 64 processes
- *Dynamical core*: e.g., density, temperature, pressure, and winds in the atmosphere (DYN group)
- *Physical parameterization*: e.g., clouds, rain, and radiation (PHYS group)

## 03\_wrf\_deimos



- Problem: 25% MPI run time fraction during the iterations of the model
- Behaviour is constant throughout all iterations
- Question: Which user function causes the problem? And why?

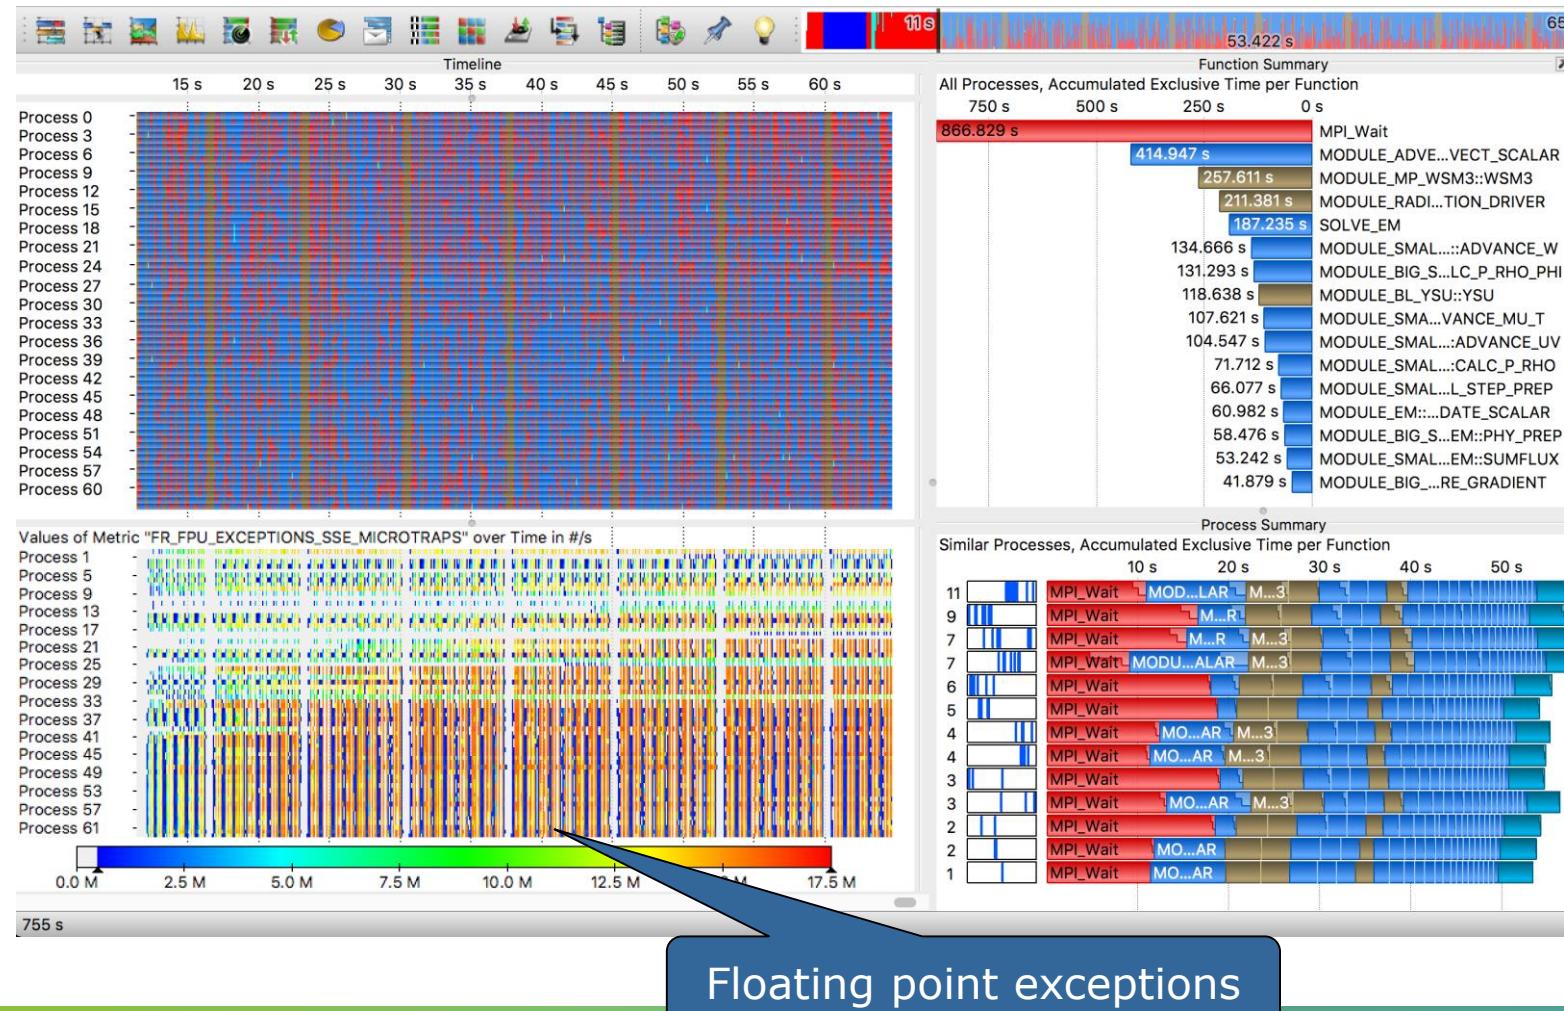
## 03\_wrf\_deimos



- Most time is spent in **MPI\_Wait**
- Top processes spent more time in **MPI\_Wait** than bottom processes
- Load imbalance in DYN group

MODULE\_ADV...VECT\_SCALAR  
runtime increases in bottom processes

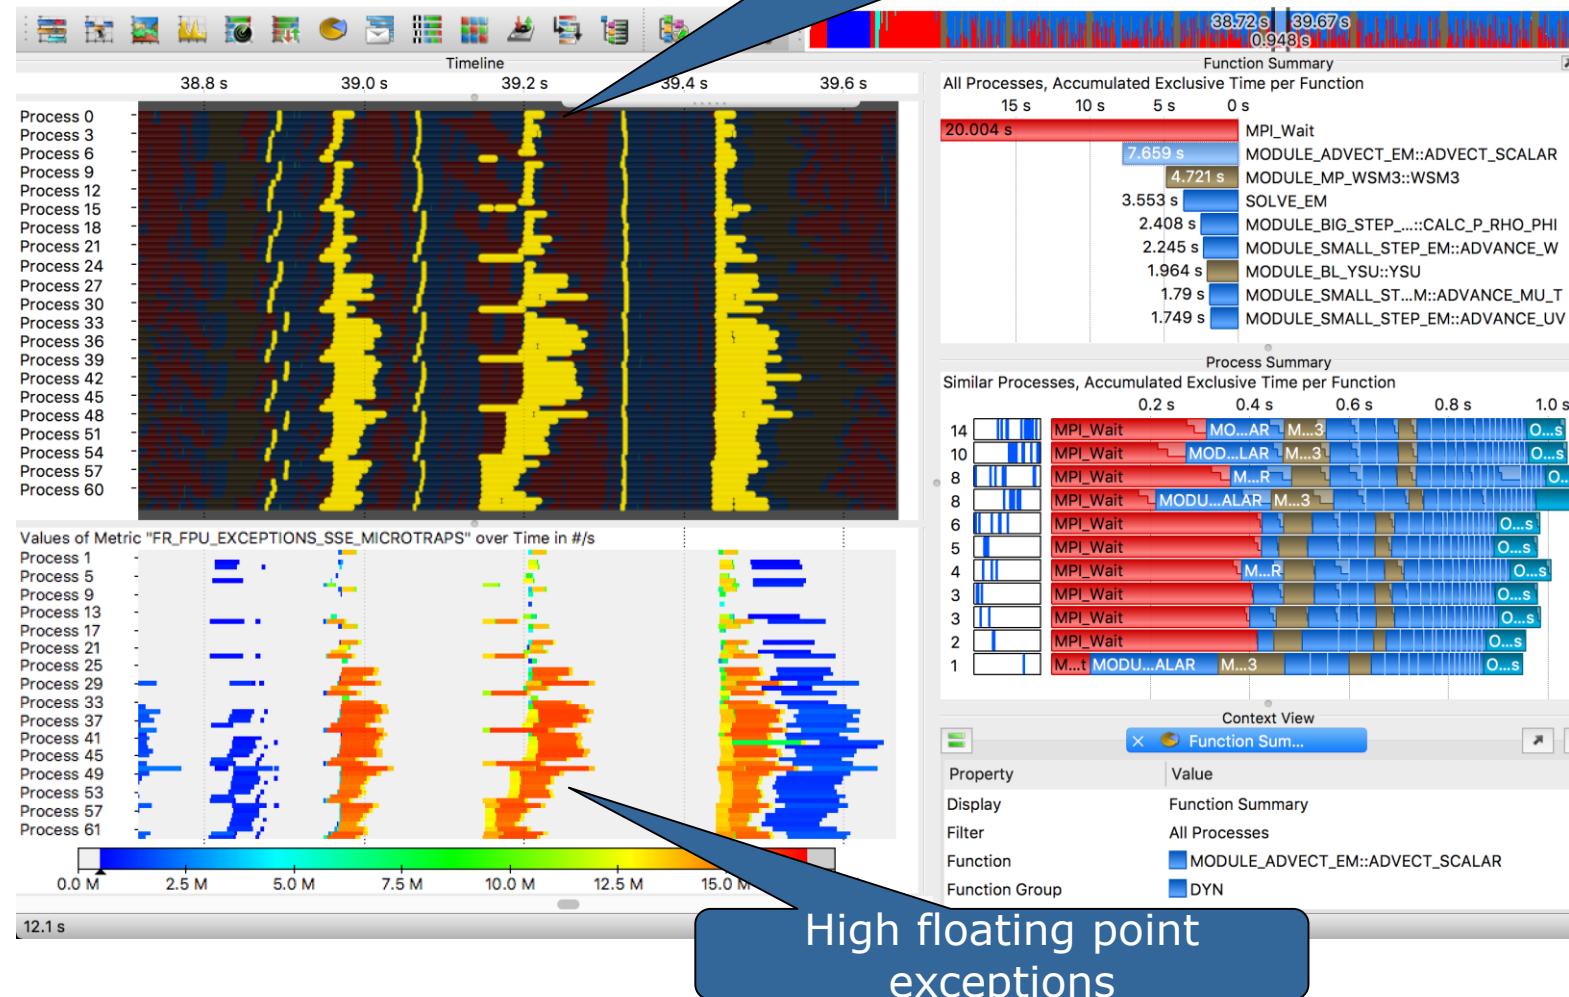
## 03\_wrf\_deimos



- Load imbalance is caused by floating point (FP) exceptions in WRF
- Counter **FR\_FPU\_EXCEPTIONS\_SSE\_MICROTRAPS** shows FP exceptions

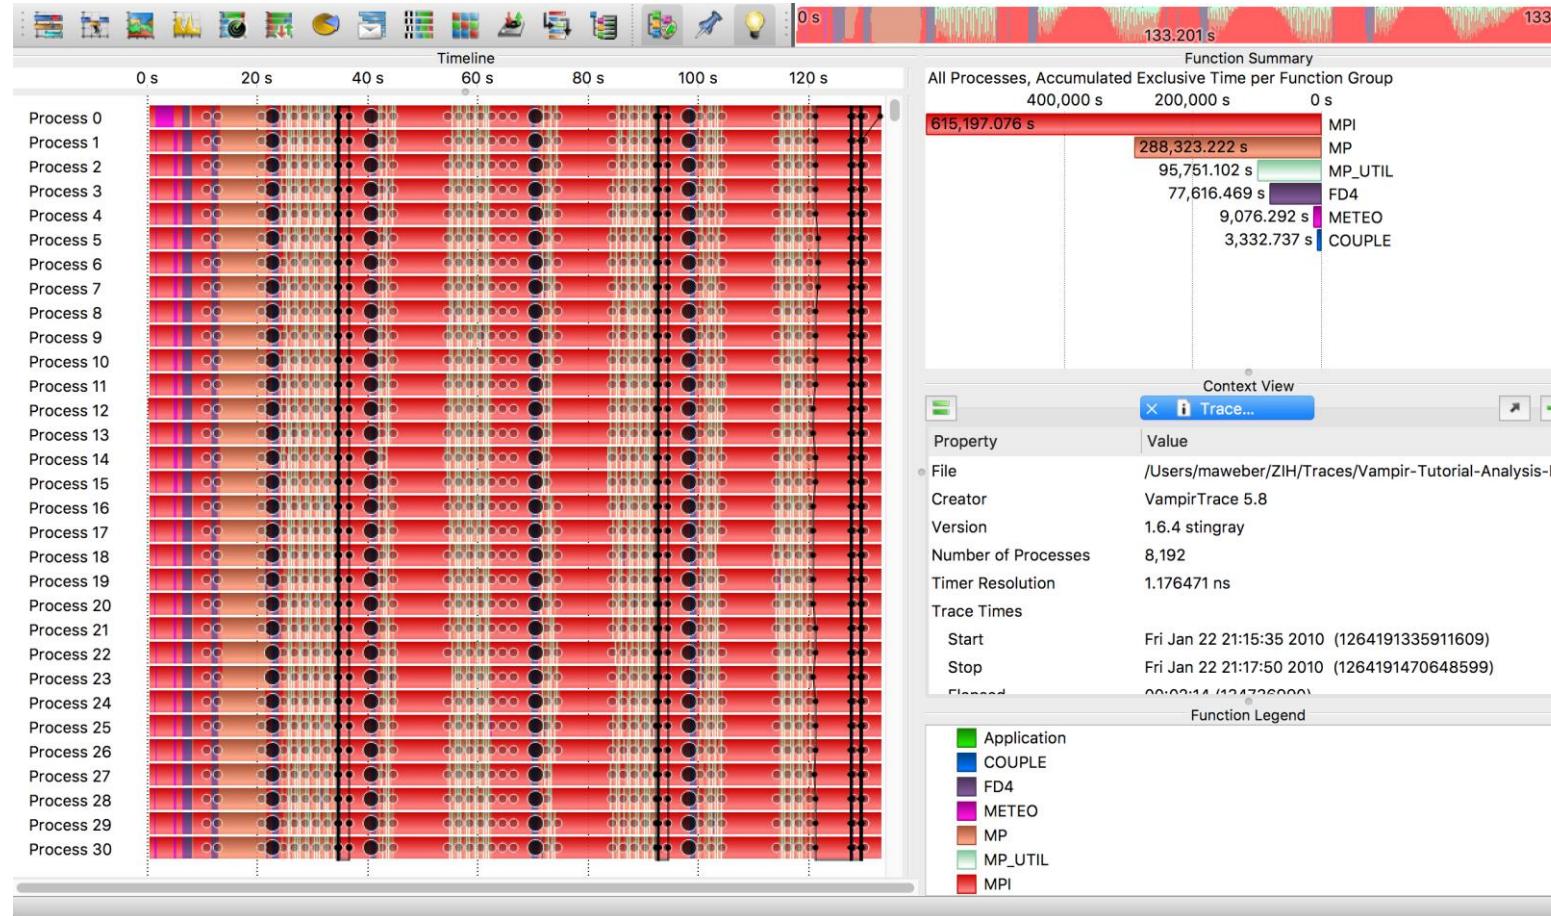
## 03\_wrf\_deimos

MODULE\_ADVECT\_EM::ADVECT\_SCALAR  
occurrences shown in yellow



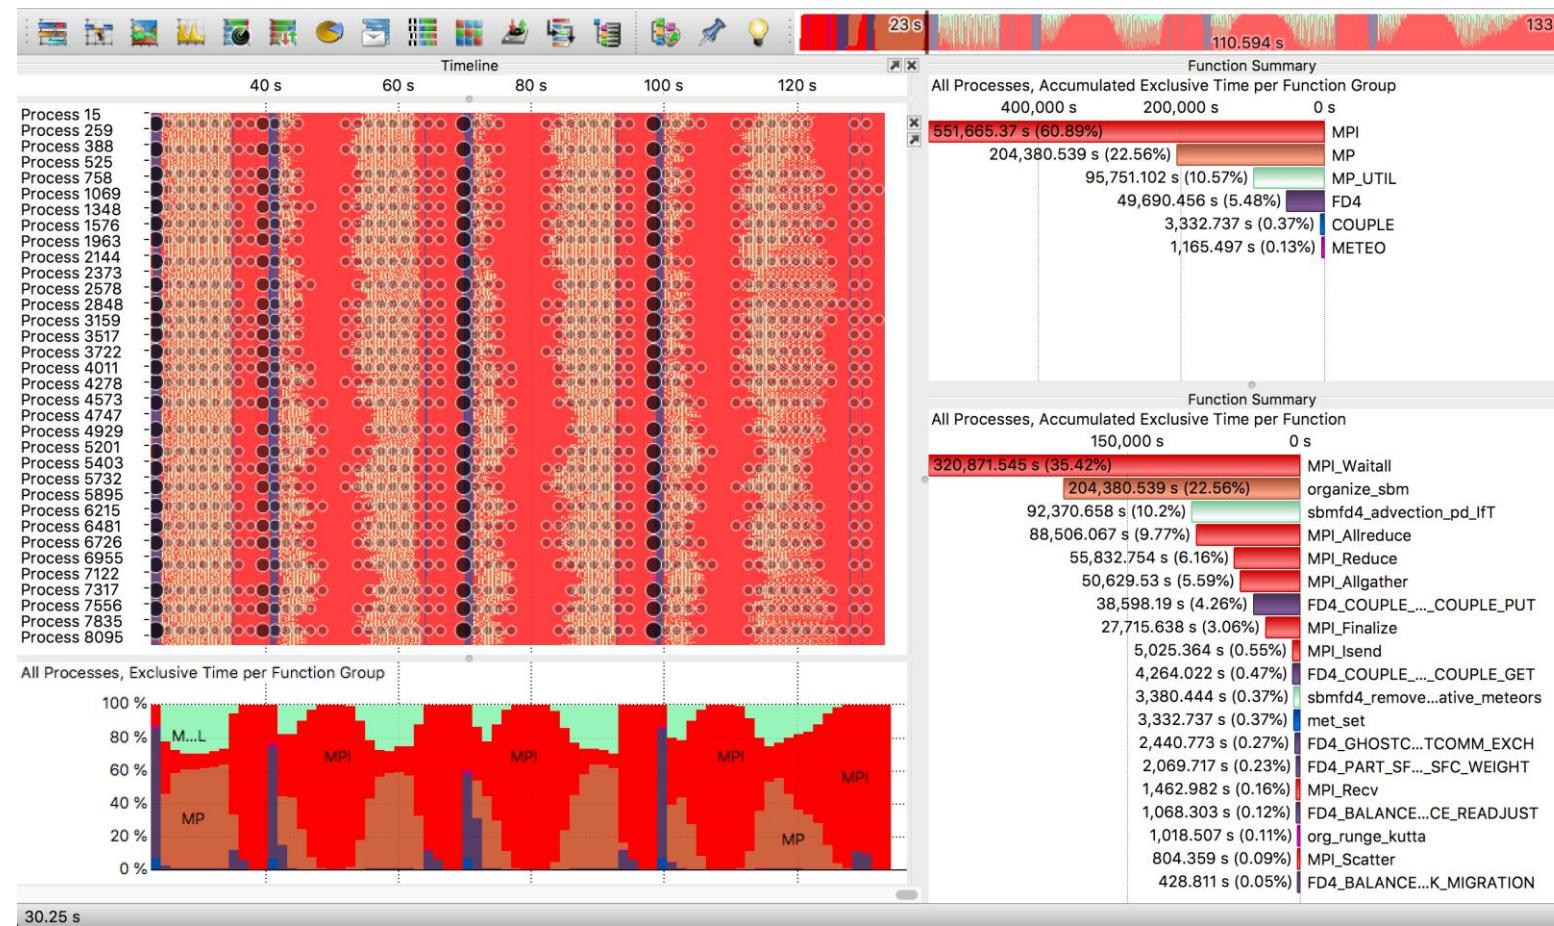
- Zoom into one iteration
- Function invocations of **MODULE\_ADVECT\_EM::ADVECT\_SCALAR** (shown in yellow) match high floating point exception occurrences indicated by the counter at the bottom

## 04\_sbmf4\_jugene



- Weather forecast code COSMO-SPECS
- Run with 8192 processes
- COSMO: weather model (METEO group)
- SPECS: microphysics for accurate cloud calculation (MP and MP\_UTIL group)
- Coupling of both models done in COUPLE group
- Dynamic load balancing (FD4 group)

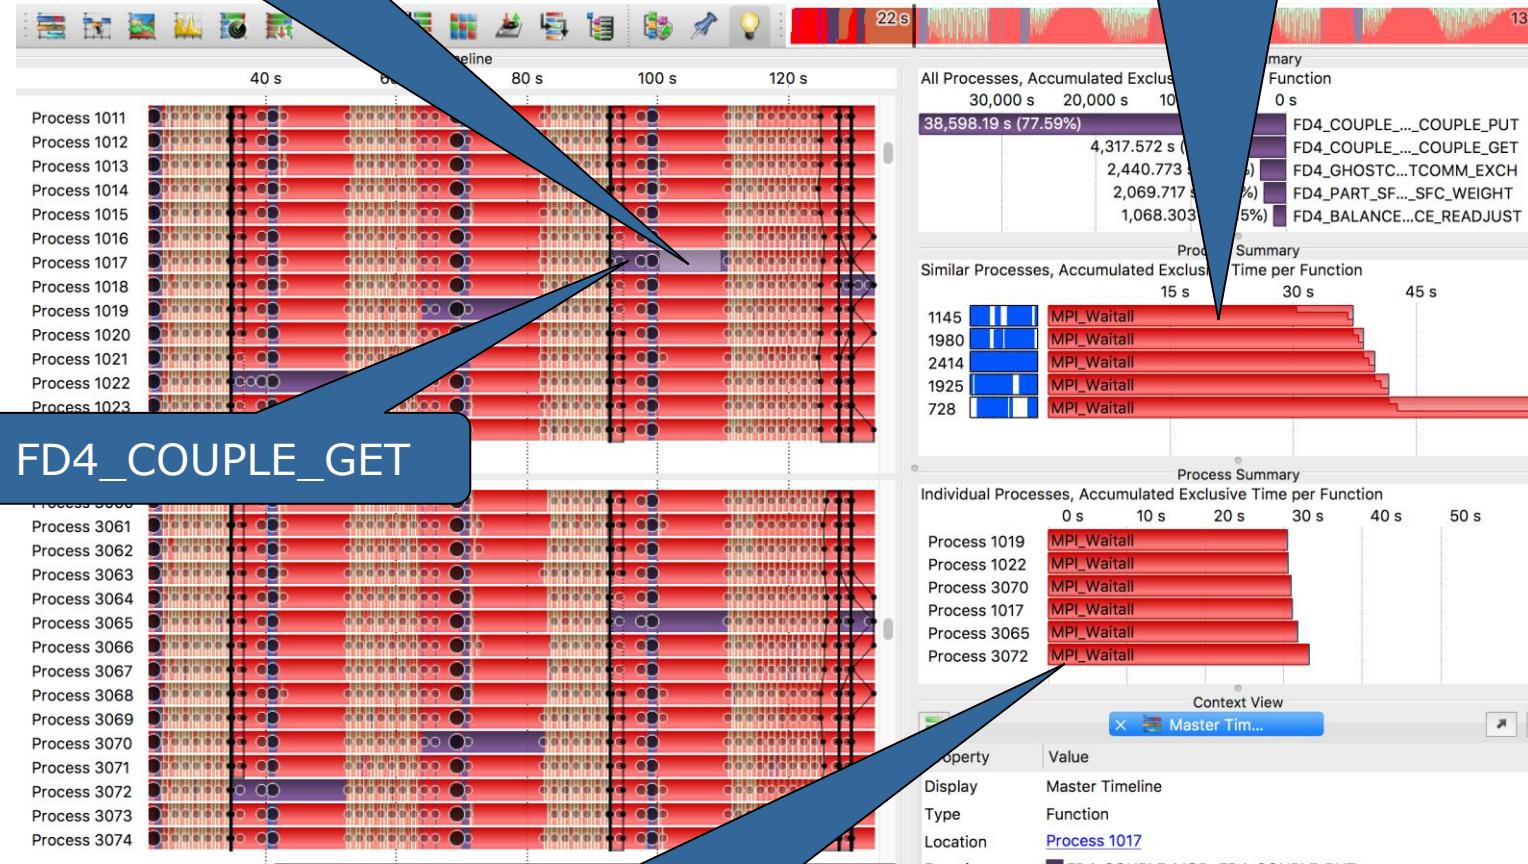
## 04\_sbmf4\_jugene



- Problem: Large MPI runtime fraction (>60%) during iterations
- Especially in **MPI\_Waitall** and **MPI\_Allreduce**
- Behaviour is constant throughout all iterations

## 04\_sbmf4\_jugene

FD4\_COUPLE\_PUT



High MPI\_Allreduce variance between processes

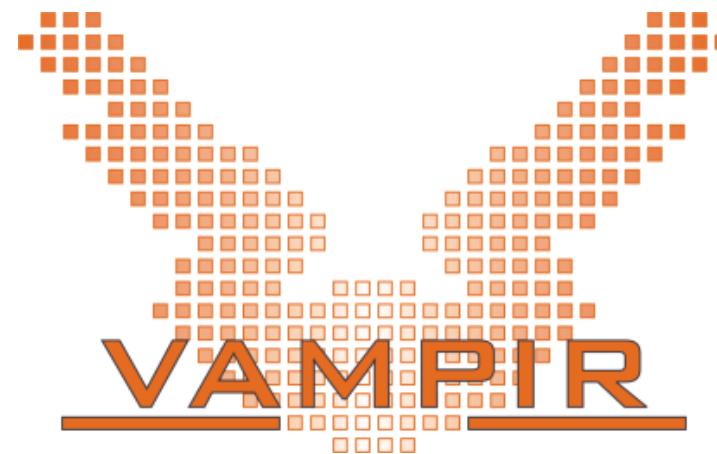
Processes sorted by  
MPI\_Allreduce timeshare

- Large runtime variation in **MPI\_Allreduce**
- Sorted profile reveals processes with small **MPI\_Waitall** timeshare
- Reason: Load imbalance in **FD4\_COUPLE\_PUT** and **FD4\_COUPLE\_GET**
- Most processes need to wait at **MPI\_Allreduce** and **MPI\_Waitall** (asynchronously)

## Summary

---

- Performance flaws can lead to significant runtime overheads
- Use resources efficiently
- Analyze your code
- Performance analysis tools are there to help you

**LMAC**Runtime Exploitation of Application Dynamism  
for Energy-efficient eXascale computing

Energy-Efficient Cluster Computing

Horizon 2020  
European Union funding  
for Research & Innovation**///H4H**    **||ParMA**

INFORMATION TECHNOLOGY FOR EUROPEAN ADVANCEMENT

<http://www.vampir.eu>[vampirsupport@zih.tu-dresden.de](mailto:vampirsupport@zih.tu-dresden.de)